

**Certainly! Here's a step-by-step guide to creating a project to predict alcohol consumption by country using Python and machine learning models. In this example, we'll use a Linear Regression model, but you can experiment with other models as well.**

#### **#### Step 1: Data Collection and Preprocessing**

**1.1. Obtain a dataset that contains alcohol consumption data by country along with relevant features (e.g., GDP, healthcare spending, population).**

**1.2. Load and preprocess the data as needed. Make sure to handle missing values and encode categorical variables if necessary.**

```
```python

import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Load the dataset
data = pd.read_csv('alcohol_consumption_data.csv')

# Drop rows with missing values
data.dropna(inplace=True)

# Select relevant features
X = data[['GDP', 'healthcare_spending', 'population']]
y = data['Total Alcohol Consumed']

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
'''
```

#### #### Step 2: Choose and Train a Machine Learning Model

**2.1. Select a machine learning model for regression. In this example, we'll use Linear Regression.**

```
```python
from sklearn.linear_model import LinearRegression

# Create a Linear Regression model
model = LinearRegression()

# Train the model
model.fit(X_train, y_train)
'''
```

#### #### Step 3: Evaluate Model Performance

**3.1. Assess the model's performance using appropriate regression metrics such as Mean Absolute Error (MAE), Mean Squared Error (MSE), and R-squared (R2).**

```
```python
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Make predictions
y_pred = model.predict(X_test)

# Calculate evaluation metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f"Mean Absolute Error: {mae:.2f}")
```

```
print(f"Mean Squared Error: {mse:.2f}")
```

```
print(f"R-squared: {r2:.2f}")
```

```
'''
```

#### #### Step 4: Interpret Model Results

**4.1. Interpret the model coefficients to understand the impact of each feature on alcohol consumption.**

```
```python
```

```
# Get the model coefficients (weights)
```

```
coefficients = model.coef_
```

```
feature_names = ['GDP', 'Healthcare Spending', 'Population']
```

```
for feature, coef in zip(feature_names, coefficients):
```

```
    print(f"{feature}: {coef:.2f}")
```

```
'''
```

#### #### Step 5: Predict Alcohol Consumption

**5.1. You can now use the trained model to predict alcohol consumption for new data by providing the relevant features (GDP, healthcare spending, population).**

```
```python
```

```
# Example prediction for a new data point
```

```
new_data_point = [[10000, 500, 5000000]] # Replace with your own values
```

```
scaled_new_data_point = scaler.transform(new_data_point)
```

```
predicted_alcohol_consumption = model.predict(scaled_new_data_point)
```

```
print(f"Predicted Alcohol Consumption: {predicted_alcohol_consumption[0]:.2f} liters")  
'''
```

#### **#### Step 6: Documentation and Reporting**

**6.1. Document your project, including data sources, preprocessing steps, model selection, evaluation metrics, and any insights gained from the analysis.**

**This example demonstrates a simple project for predicting alcohol consumption by country using a Linear Regression model. Depending on the dataset and the complexity of the problem, you can explore more advanced models and techniques for better predictions.**